

Web-scale Knowledge Acquisition

Author: Eric Nichols, eric@ecei.tohoku.ac.jp

Web-KA: A collection of tools for acquiring and databasing semantic relation patterns and instances from Web texts. Contains Python implementations of the Espresso[1,2] and Coupled Pattern Learner [3] bootstrapping algorithms and tools for managing patterns and instances in MongoDB.

instances2matrix.py

`instances2matrix.py`: creates a matrix of co-occurrence counts between relation pattern * arguments in mongodb from input instances

Usage

```
Usage: instances2matrix.py [options] [<instance_files>]
```

Options:

```
-h, --help            show this help message and exit
-c COLLECTION, --collection=COLLECTION
                        collection name
-d DB, --database=DB  database name
-o HOST, --host=HOST  mongodb host machine name.
                        default: localhost
-p PORT, --port=PORT  mongodb host machine port number.
                        default: 27017
```

Instances

Format

Instances have the following tab-delimited format:

- `score`: score representing weight * co-occurrence count for instance
- `loc`: giving source and location of instance
- `rel`: containing relation pattern
- `argc`: giving argument count
- `argv`: tab-delimited list of arguments as strings

Example

```
1.0\treverb_clueweb_tuples-1.1.txt:30:10-11\tARG1 acquired  
ARG2\t2\Google\tYouTube
```

Co-occurrence Matrix

Format

The co-occurrence matrix collection has the following fields:

- `rel`: relation pattern
- `arg1`: first argument
- ...
- `argn`: nth argument
- `score`: score for $\text{rel} * \text{args tuple}$

Naming Scheme

Instances of differing argument count are stored in separate mongodb collections with names formatted as `<collection>_<argc>`. E.g. if a collection `clueweb` has instances with argument counts of 1, 2, and 3, then the following collection would be created:

- `clueweb_1`
- `clueweb_2`

- `clueweb_3`

Indexing

It is indexed for fast look up of rel, args, and (rel,args) tuples.

matrix2pmi.py

`matrix2pmi.py`: caches co-occurrence frequencies and discounted PMI between relation patterns and argument tuples into a matrix stored in mongodb

Usage

```
Usage: matrix2pmi.py [options] [database] [collection]
```

Options:

```
-h, --help            show this help message and exit
-o HOST, --host=HOST  mongodb host machine name.
                       default: localhost
```

```
-p PORT, --port=PORT  mongodb host machine port number.  
                        default: 1979  
  
-s START, --start=START  
                        specify calculation to start with  
                        1 or F_i: instance tuple frequencies  
                        2 or F_p: relation pattern frequencies  
                        3 or F_ip: instance*pattern co-occurrence  
frequencies  
  
                        4 or pmi_ip: instance*pattern discounted  
PMI score  
  
                        default: F_i
```

Caches Created

Creates 4 frequency/score caches in the form of mongodb collections:

1. `<matrix>_F_i`: instance tuple frequencies

2. `<matrix>_F_p`: relation pattern frequencies
3. `<matrix>_F_ip`: instance*pattern co-occurrence frequencies
4. `<matrix>_pmi_ip`: instance*pattern Pointwise Mutual Information score discounted to account for bias toward infrequent events following [1]

Pointwise Mutual Information

Pointwise mutual information between argument instances and relation patterns is defined following [2] as:

$$(1) \text{ PMI}(i, p) = \log(F(i, p) / F(i) * F(p))$$

where

- (2) $F(i)$ = the frequency of argument instance i
- (3) $F(p)$ = the frequency of relation pattern p
- (4) $F(i,p)$ = the co-occurrence frequency of argument instance i and relation pattern p

Discounted PMI

Pointwise Mutual Information is known to be biased toward infrequent events. Pantel and Ravichandran [1] compensate by multiplying PMI by a “discounting factor” that is essentially a smoothed co-occurrence frequency multiplied by a smoothed frequency of the argument instance or the relation pattern, whichever is lesser.

- (5) $\text{discount}(i,p) = (F(i,p) / (F(i,p)+1)) * (\min(F(i),F(p)) / (\min(F(i),F(p))+1))$
- (6) $\text{discountedPMI}(i,p) = \text{PMI}(i,p) * \text{discount}(i,p)$

espresso.py

`espresso.py`: an implemenatation of the Espresso bootstrapping algorithm

Usage

```
Usage: espresso.py [options] [database] [collection] [rel]
[seeds]
```

Options:

```
-h, --help            show this help message and exit
-o HOST, --host=HOST  mongodb host machine name. default:
localhost
-p PORT, --port=PORT  mongodb host machine port number.
default: 27017
-s START, --start=START
                        iteration to start with. default: 1
-t STOP, --stop=STOP  iteration to stop at. default: 2
```

Caches Created

Creates 2 caches of bootstrapped instances and patterns for the target relation:

1. `<matrix>_<rel>_esp_i`: bootstrapped instances for

2. `<matrix>_<rel>_esp_p`: bootstrapped patterns for

Bootstrapping

Bootstrapping starts with seed instances and alternates between promoting new patterns and instances following the Espresso bootstrapping algorithm [2].

1. retrieve promoted instances/patterns
2. rank by reliability score
3. keep top 10 promoted instances/patterns
4. bootstrap patterns/instances using promoted instances/patterns

Reliability Score

Candidate patterns and instances are ranked by reliability score, which reflects the pointwise mutual information score between a promoted pattern/instance and the set of instances/patterns that generated it.

$$(1) \quad r_i(i, P) = \frac{\sum_{p \in P} \text{dpmi}(i, p) * r_p(p)}{\max_{p \in P} \text{dpmi}(i, p)} \quad \text{for } p \text{ in } P$$

$$(2) \quad r_p(P, i) = \frac{\sum_{i \in I} \text{dpmi}(i, p) * r_i(i)}{\max_{i \in I} \text{dpmi}(i, p)} \quad \text{for } i \text{ in } I$$

where dpmi is Discounted Pointwise Mutual Information [1]. r_i and r_p are recursively defined with $r_i=1.0$ for the seed instances.

References

- [1] Patrick Pantel and Deepak Ravichandran. Automatically Labeling Semantic Classes. HLT-NAACL 2004.
- [2] Patrick Pantel and Marco Pennacchiotti. Espresso: Leveraging Generic Patterns for Automatically Harvesting Semantic Relations. ACL 2006.
- [3] Andrew Carlson, Justin Betteridge, Richard C. Wang, Estevam R. Hruschka Jr.,

and Tom M. Mitchell. Coupled Semi-supervised Learning for Information Extraction. WDSM 2010.